



Databases

Basic SQL



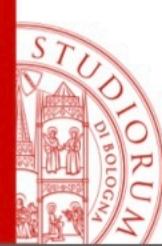
SQL

- At first, an acronym for "Structured Query Language", now a “proper noun”
- Language with many features:
 - Implements both DDL and DML
- There is an standard ISO language, but different DBMSs have their own language grammar.
- For the moment we’re going to see the basics of this language



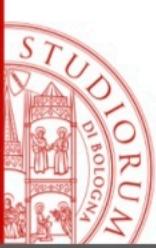
SQL: History

- Its predecessor was **SEQUEL** (1974);
- First implementations were SQL/DS and Oracle (1981)
- SQL has a “de facto” standard since 1983
- Many proposed updates (1986, then 1989, 1992, 1999, 2003, 2006, 2008, ...) but still, DBMSs have their own grammar (!! See some comparisons on
<http://troels.arvin.dk/db/rdbms/>)



SQL improvements: SQL-base

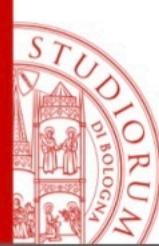
- **SQL-86:** first proposed standard. It had most of the clauses for expressing queries, but offered a limited support for creating and updating both schemas and data.
- **SQL-89:** Referential Integrity is added



SQL improvements: SQL-2

SQL-92: mostly backward compatible, has new features:

- New functions (e.g. COALESCE, NULLIF, CASE)
- 3 usage levels: entry, intermediate, full

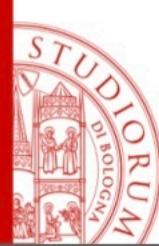


SQL Improvements: SQL-3 (1)

Different subversions:

- **SQL:1999:** proposes the object-relational, triggers and extern functions

- **SQL:2003:** extends the object oriented model and allows to perform queries in Java and over semistructured data (XML)



SQL Improvements: SQL-3 (2)

- **SQL:2006:** SQL is extended with other languages (e.g. XQuery) for querying XML data
- **SQL:2008:** some slight edits to the syntax (e.g. trigger with instead of)



SQL Improvements

Unofficial Name	Official Name	Features
SQL-Base	SQL-86	Basic keywords
	SQL-89	Referential Integrity
SQL-2	SQL-92	Modello relazionale New keywords 3 levels: entry, intermediate, full
SQL-3	SQL:1999	Relational model with object-oriented Structured in different parts Trigger, external functions, ...
	SQL:2003	The support of the Object-Oriented model is extended The no-longer used keywords were removed Extensions: SQL/JRT, SQL/XML, ...
	SQL:2006	Extended support for XML data
	SQL:2008	Slight edits (e.g.: trigger instead of)



Data Definition (1)

CREATE DATABASE:

- Each newly created database contains tables, views, triggers and other things

E.g.:

CREATE DATABASE db_name

please note:

In **SQLite** `sqlite3_open_v2(db_name)`

In **Mimer** `CREATE DATABANK db_name`



Data Definition (2)

CREATE SCHEMA:

- A SQL Schema is identified by a name and describes the elements belonging to it (tables, types, constraints, views, domains, ...). The schema will belong to the user which has typed the statement

Example:

CREATE SCHEMA schema_name



Data Definition (3)

CREATE SCHEMA:

- Such statement could be even followed by the **AUTHORIZATION** keyword, to indicate a specific user owning the schema.

Example:

```
CREATE SCHEMA schema_name
AUTHORIZATION 'user_name'
```



Data Definition (4)

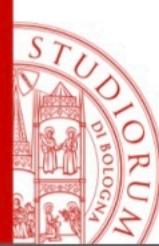
CREATE TABLE:

- Specifies a new relation and creates its empty instance.
- It specifies its attributes (with their types) and initial constraints.



CREATE TABLE: example

```
CREATE TABLE Employee(  
    Number CHAR(6) PRIMARY KEY,  
    Name CHAR(20) NOT NULL,  
    Surname CHAR(20) NOT NULL,  
    Dept CHAR(15),  
    Wage NUMERIC(9) DEFAULT 0,  
    FOREIGN KEY(Dept) REFERENCES  
        Department(Dept),  
    UNIQUE (Surname,Name)  
)
```



(Attribute) Data Types

- (Attribute) Data Types in SQL correspond to the domains in the relational calculus.
- *Basic* data types (already available)
- *Custom* data types (called “domains” simple and reusable)



Basic Data Types

- **Character-string**: data types are either fixed length or varying length.
- **Numeric**, including integer numbers and different floating points.
- **DATE, TIME, INTERVAL**
- Introduced with SQL-3:
 - **Boolean**
 - **BLOB, CLOB** (binary/character large object): representing huge data collections (either textual or not)



Custom Data Types

CREATE DOMAIN:

- Each custom data type could be used when defining new relations, stating constraints and default values



CREATE DOMAIN: example

```
CREATE DOMAIN Grade
AS SMALLINT DEFAULT NULL
CHECK ( value >=18 AND value <= 30 )
```



Table Constraints

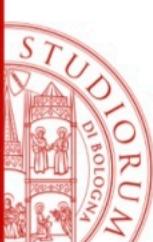
- **NOT NULL**
- **UNIQUE** defining keys
- **PRIMARY KEY**: (just one, implies **NOT NULL**;
DB2 has a non standard behaviour)
- **CHECK**, let's see it later



UNIQUE and PRIMARY KEY

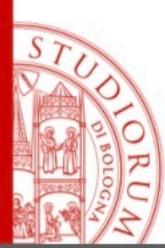
It could be used when:

- when we define an attribute that defines the key.
- as a stand-alone element



CREATE TABLE: example

```
CREATE TABLE Employee(  
    Number CHAR(6) PRIMARY KEY,  
    Name CHAR(20) NOT NULL,  
    Surname CHAR(20) NOT NULL,  
    Dept CHAR(15),  
    Wage NUMERIC(9) DEFAULT 0,  
    FOREIGN KEY(Dept) REFERENCES  
        Department(Dept),  
    UNIQUE (Surname,Name)  
)
```



PRIMARY KEY, other options

Number **CHAR(6)** **PRIMARY KEY**

Number **CHAR(6)**,

PRIMARY KEY (Number)



CREATE TABLE: example

```
CREATE TABLE Employee(  
    Number CHAR(6) PRIMARY KEY,  
    Name CHAR(20) NOT NULL,  
    Surname CHAR(20) NOT NULL,  
    Dept CHAR(15),  
    Wage NUMERIC(9) DEFAULT 0,  
    FOREIGN KEY(Dept) REFERENCES  
        Department(Dept),  
    UNIQUE (Surname,Name)  
)
```



Warning!

Name **CHAR(20) NOT NULL,**
Surname **CHAR(20) NOT NULL,**
UNIQUE (Surname,Name)

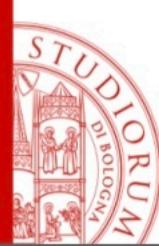
Name **CHAR(20) NOT NULL UNIQUE,**
Surname **CHAR(20) NOT NULL UNIQUE,**

- Are not the same thing!



Key and Referential Integrity Constraints

- **CHECK**, let's see it later
- **REFERENCES** and **FOREIGN KEY** define Referential Integrity Constraints
- They can be defined
 - Over a single attribute
 - Over multiple attributes
- We can define *referential triggered actions* when such constraints are violated



Referential Integrity Constraints (1)

Offenses	<u>Code</u>	Day	Officer	State	Number
	34321	1/2/95	3987	IT	AG548UK
	53524	4/3/95	3295	IT	TE395AB
	64521	5/4/96	3295	FR	ZT395AB
	73321	5/2/98	9345	FR	ZT395AB

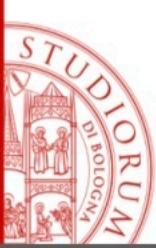
Officer	<u>Number</u>	Surname	Name		
	3987	Rossi	Luca		
	3295	Neri	Piero		
	9345	Neri	Mario		
	7543	Mori	Gino		



Referential Integrity Constraints (2)

Offenses	<u>Code</u>	Day	Officer	State	Number
	34321	1/2/95	3987	IT	AG548UK
	53524	4/3/95	3295	IT	TE395AB
	64521	5/4/96	3295	FR	ZT395AB
	73321	5/2/98	9345	FR	ZT395AB

Car	<u>State</u>	<u>Number</u>	Surname	Name
	IT	AG548UK	Verdi	Giuseppe
	IT	TE395AB	Verdi	Giuseppe
	FR	ZT395AB	Quinault	Philippe



CREATE TABLE: example

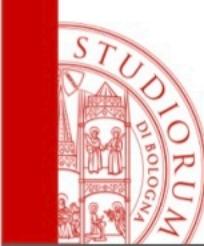
```
CREATE TABLE Offences(
    Code CHAR(6) NOT NULL PRIMARY KEY,
    Day DATE NOT NULL,
    Officer INTEGER NOT NULL
        REFERENCES Officer(Number),
    State CHAR(2),
    Number CHAR(6),
    FOREIGN KEY(State, Number)
        REFERENCES Car(State, Number)
)
```



Referential Triggered Action

- After each referential constraint, we can specify a triggered action (delete, update) to be invoked if the operation is rejected:

```
ON < DELETE | UPDATE >
    < CASCADE | SET NULL |
        SET DEFAULT | NO ACTION >
```



Referential Triggered Action: delete

- **CASCADE**: deletes the referencing tuples.
- **SET NULL**: the value of the deleted referencing attribute is replaced with NULL.
- **SET DEFAULT**: the value of the deleted referencing attributes is replaced with the specified default value.
- **NO ACTION**: no removal is allowed



Referential Triggered Action: update

- **CASCADE**: the value of the referencing foreign key attribute(s) is updated with the new value.
- **SET NULL**: the value of the affected referencing attribute is replaced with NULL.
- **SET DEFAULT**: the value of the affected referencing attributes is replaced with the specified default value.
- **NO ACTION**: no update is allowed.



Schema Change Statements

■ **ALTER DOMAIN**

■ **ALTER TABLE**

■ **DROP DOMAIN**

■ **DROP TABLE**



ALTER DOMAIN

ALTER DOMAIN:

- Allows to alter previously-defined domains.
- Such statement has to be used alongside with those other ones: **SET DEFAULT**, **DROP DEFAULT**, **ADD CONSTRAINT** or **DROP CONSTRAINT**



ALTER DOMAIN: example 1

ALTER DOMAIN Grade SET DEFAULT 30

- Sets the default **Grade** to 30
- Such command is applied only when the command is invoked and missing grade value are found

ALTER DOMAIN Grade DROP DEFAULT

- Removes the default **Grade** value



ALTER DOMAIN: example 2

ALTER DOMAIN Grade

SET CONSTRAINT isValid **CHECK** (value
 ≥ 18 **AND** value ≤ 30)

- Adds the isValid constraint to the data type Grade

ALTER DOMAIN Grade

DROP CONSTRAINT isValid

- Removes constraint associated to the data type



ALTER TABLE

ALTER TABLE:

- Performs changes to previously defined tables
- Such statement has to be used alongside with these parameters: **ALTER COLUMN**, **ADD COLUMN**, **DROP COLUMN**, **DROP CONSTRAINT** or **ADD CONSTRAINT**



ALTER TABLE: example 1

ALTER TABLE Employee

ALTER COLUMN Number SET NOT NULL

- **Number** from table **Employee** cannot have null values

ALTER TABLE Employee

ADD COLUMN Level CHARACTER(10)

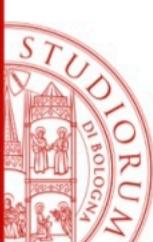
- An attribute **Level** is added to the table **Employee**.



How to interpret the SELECT clause

```
3 SELECT Number, Name  
1 FROM Officer  
2 WHERE Surname = 'Jones'
```

- 1 From the table **Officer**
- 2 Retrieve all the officers having
'Jones' as **Surname** attribute
- 3 Showing for each tuple both
Number and **Name**



ALTER TABLE: example 2

ALTER TABLE Employee

DROP COLUMN Level RESTRICT

- Removes the attribute **Level** from **Employee** only if it doesn't contain values

ALTER TABLE Employee

DROP COLUMN Level CASCADE

- Removes the attribute **Level** from **Employee** alongside with its values



ALTER TABLE: example 3

ALTER TABLE Employee

ADD CONSTRAINT validNum CHECK
(char_length(Number) = 10)

- Adds the validNum constraint to the Number attribute from Employee

ALTER TABLE

DROP CONSTRAINT

Employee

validNum

- Removes the previously defined constraint



DROP DOMAIN

DROP DOMAIN:

- Removes a user-defined data type

Example:

DROP DOMAIN Grade



DROP TABLE

DROP TABLE:

- Removes a whole table instance with its schema and its data

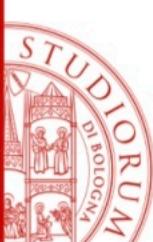
Example:

DROP TABLE Offences



Defining Indices

- They usually enhance the query time, relevant for computation efficiency
- They are defined at the physical level, not logical
- In the old days this was also the only way to define keys
- **CREATE INDEX**



CREATE INDEX: example

```
CREATE INDEX idx_Surname  
ON Officer (Surname)
```

- Creates the index `idx_Surname` on the attribute `Surname` from the table `Officer`



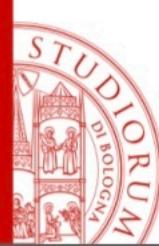
DDL in practice

- In many systems and projects, different tools are used, instead of SQL statements, in order to define a database schema (e.g. tools with a graphical user interface).



SQL, data operations

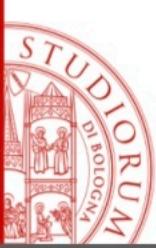
- Query:
 - **SELECT**
- Edit:
 - **INSERT, DELETE, UPDATE**



SELECT, shortcuts (1)

```
SELECT *
FROM People
WHERE Age < 30
```

```
SELECT Name, Age, Income
FROM People
WHERE Age < 30
```



Basic SELECT statement

```
SELECT <AttributeList>
FROM <TableList>
[ WHERE <Condition> ]
```

- Target list
- **FROM** statement
- **WHERE** statement



Database Example 2

People

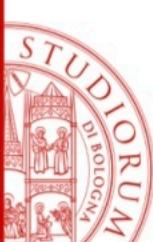
Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Jesse	50	35
Phil	26	30
Louis	50	40
Frank	60	20
Olga	30	41
Steve	85	35
Abby	75	87

Motherhood

Mother	Child
Abby	Alice
Abby	Louis
Jesse	Olga
Jesse	Phil
Alice	Jim
Alice	James

Fatherhood

Father	Child
Steve	Frank
Louis	Olga
Louis	Phil
Frank	Jim
Frank	James



Selection and Projection

Return name and income of people being less than 30 yo.

$$\pi_{\text{Name, Income}}(\sigma_{\text{Age} < 30}(\text{People}))$$

```
SELECT name, income  
FROM people  
WHERE age < 30
```

People

Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Jesse	50	35
Phil	26	30
Louis	50	40
Frank	60	20
Olga	30	41
Steve	85	35
Abby	75	87



Selection and Projection

Return name and income
of people being less than
30 yo.

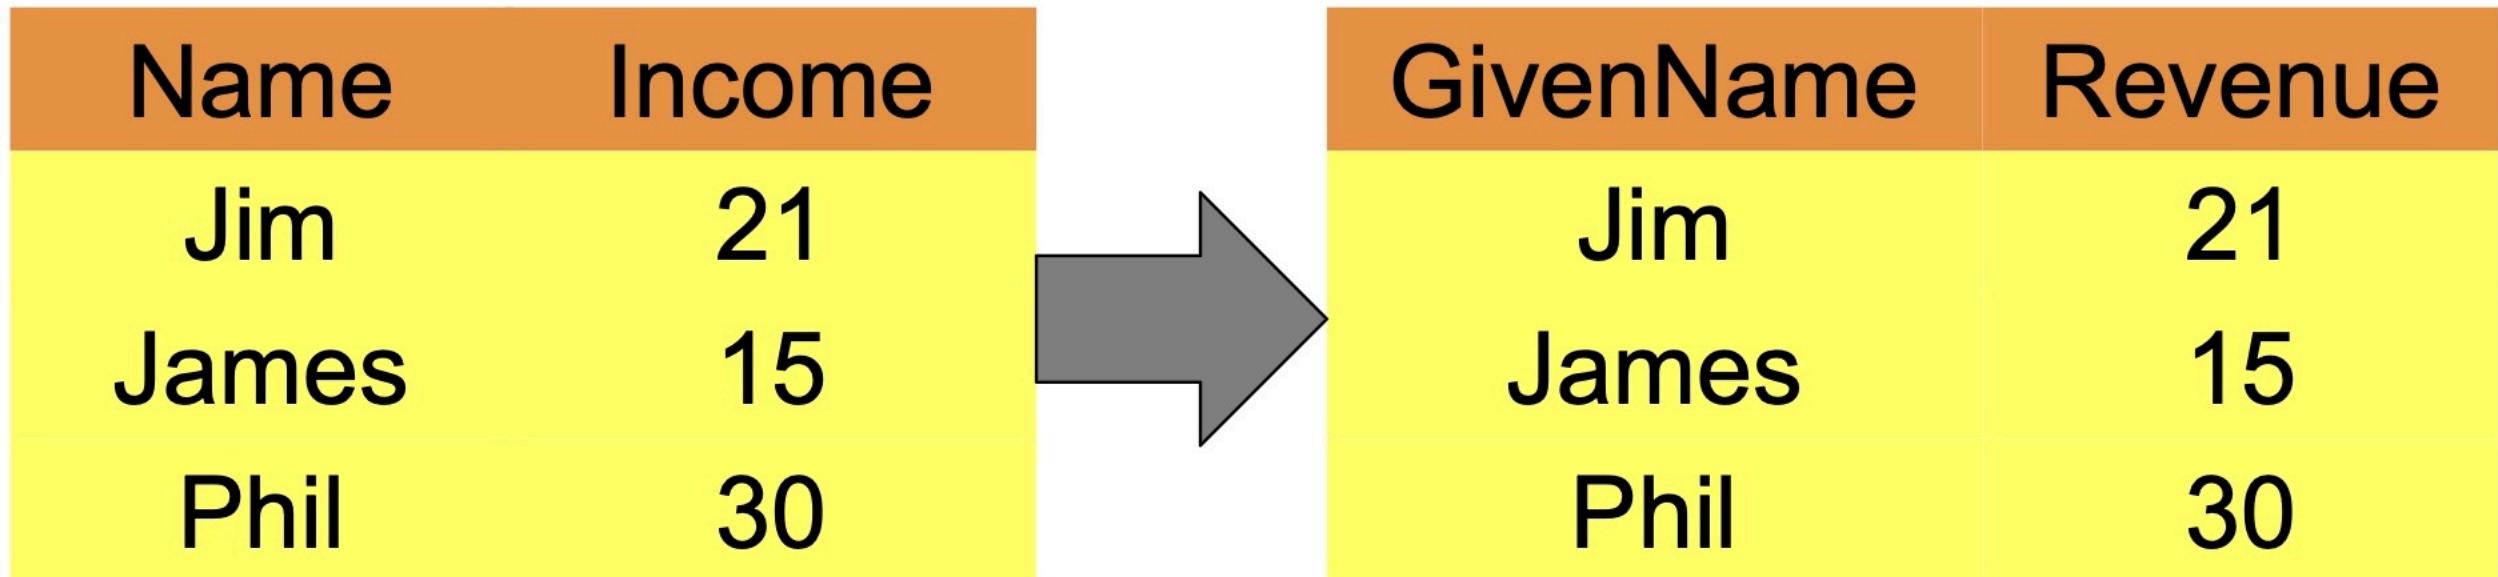
$$\pi_{\text{Name, Income}}(\sigma_{\text{Age} < 30}(\text{People}))$$

Name	Income
Jim	21
James	15
Phil	30

```
SELECT Name, Income  
FROM People  
WHERE age < 30
```

SELECT, (attribute) renaming

```
SELECT p.Name AS GivenName,
      p.Income AS Revenue
FROM People AS p
WHERE p.Age < 30
```



The diagram illustrates the transformation of a table using SQL's SELECT statement with column renaming. On the left, a table has columns 'Name' and 'Income'. An arrow points to the right, where the same data is shown with renamed columns: 'GivenName' and 'Revenue'.

Name	Income	GivenName	Revenue
Jim	21	Jim	21
James	15	James	15
Phil	30	Phil	30



Pure Selection

Provide the Name, Age and Income of the people begin less than 30 yo.

$\sigma_{Age < 30}(\text{People})$

```
SELECT *
FROM People
WHERE Age < 30
```

People			
Name	Age	Income	
Jim	27	21	
James	25	15	
Alice	55	42	
Jesse	50	35	
Phil	26	30	
Louis	50	40	
Frank	60	20	
Olga	30	41	
Steve	85	35	
Abby	75	87	



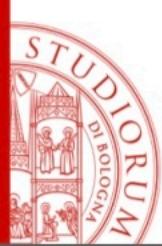
Selection, without projection

Provide the Name, Age and Income of the people begin less than 30 yo.

$\sigma_{Age < 30}(\text{People})$

Name	Age	Income
Jim	27	21
James	25	15
Phil	26	30

```
SELECT *
FROM People
WHERE Age < 30
```



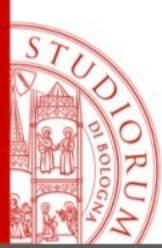
Projection, without selection

Return the peoples' name and income

$$\pi_{\text{Name}, \text{Income}}(\text{People})$$

```
SELECT Name, Income  
FROM People
```

People			
Nome	Age	Income	
Jim	27	21	
James	25	15	
Alice	55	42	
Jesse	50	35	
Phil	26	30	
Louis	50	40	
Frank	60	20	
Olga	30	41	
Steve	85	35	
Abby	75	87	



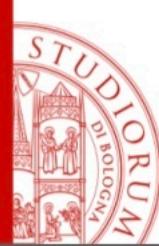
Projection, without selection

Return the peoples' name and income

$\pi_{\text{Name}, \text{Income}}(\text{People})$

SELECT Name, Income
FROM People

Name	Income
Jim	21
James	15
Alice	42
Jesse	35
Phil	30
Louis	40
Frank	20
Olga	41
Steve	35
Abby	87



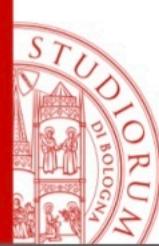
SELECT, shortcuts (2)

Given a relation R(A,B)

```
SELECT *
FROM R
```

It corresponds to:

```
SELECT X.A AS A, X.B AS B
FROM R X
WHERE true
```

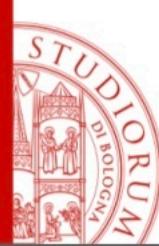


Composed Conditions

```
SELECT *
FROM People
WHERE Income > 25 AND (Age < 30 OR Age > 60)
```

People

Name	Age	Income
Phil	26	30
Frank	60	20
Olga	30	41
Steve	85	35



Composed Conditions

```
SELECT *
FROM People
WHERE Income>25 AND (Age<30 OR Age>60)
```

Name	Age	Income
Phil	26	30
Steve	85	35



“LIKE” Predicate

- It returns the people having a name starting with 'J' and have a 'm' as a third letter:

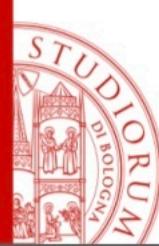
```
SELECT *
FROM People
WHERE Name LIKE 'J_m%'
```



“LIKE” Predicate

```
SELECT *
FROM People
WHERE Name LIKE 'J_m%'
```

Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Jesse	50	35
Phil	23	30
Louis	50	40
Frank	60	20



Handling NULL values

Employee

Number	Surname	Agency	Age
5998	Neri	Milan	45
9553	Bruni	Milan	NULL

- Return the employees being either more than 40 yo or NULL value

$$\sigma_{(\text{Age} > 40) \text{ OR } (\text{Age IS NULL})} (\text{Employee})$$



Example

- Return the employees being either more than 40 yo or NULL value

$$\sigma_{(Age > 40) \text{ OR } (Age \text{ IS NULL})} (\text{Employee})$$

```
SELECT *
FROM Employee
WHERE Age>40 OR Age IS NULL
```



Projection (Relational Algebra)

Employee

Number	Surname	Agency	Income
7309	Neri	Naples	55
5998	Neri	Milan	64
9553	Rossi	Rome	44
5698	Rossi	Rome	64

- Return the surname and the agency for all the employees

$$\pi_{\text{Surname}, \text{Agency}} (\text{Employee})$$



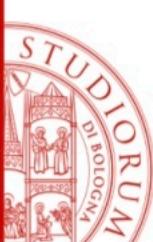
Projection (SQL and DISTINCT)

```
SELECT  
    Surname, Agency  
FROM Employee
```

Surname	Agency
Neri	Naples
Neri	Milan
Rossi	Rome
Rossi	Rome

```
SELECT DISTINCT  
    Surname, Agency  
FROM Employee
```

Surname	Agency
Neri	Naples
Neri	Milan
Rossi	Rome



Select, Project, Join

- By using only one relation in the **FROM** clause, one single SQL query can express: *select*, *project* and *rename*
- Using more relations in the **FROM** clause we have *joins* (and cartesian products)



SQL and Relational Algebra (1)

- $R_1(A_1, A_2) \ R_2(A_3, A_4)$

```
SELECT DISTINCT R1.A1, R2.A4
FROM    R1, R2
WHERE   R1.A2 = R2.A3
```

- Cartesian products (**FROM**)
- Selection (**WHERE**)
- Projection (**SELECT**)



SQL and Relational Algebra (2)

- $R1(A1, A2) \ R2(A3, A4)$

```
SELECT DISTINCT R1.A1, R2.A4
FROM      R1, R2
WHERE    R1.A2 = R2.A3
```

$$\pi_{A1, A4} (\sigma_{A2=A3} (R1 \bowtie R2))$$



SQL, alias and renaming

- renaming could be required
 - in the cartesian product
 - in the target list

```
SELECT X.A1 AS B1, ...
FROM   R1 X, R2 Y, R1 Z
WHERE X.A2 = Y.A3 AND ...
```



SQL vs. Relational Algebra

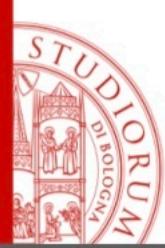
```
SELECT DISTINCT X.A1 AS B1, Y.A4 AS  
      B2  
FROM    R1 X, R2 Y, R1 Z  
WHERE   X.A2 = Y.A3 AND Y.A4 = Z.A1
```

$$\begin{aligned} & \rho_{B1,B2 \leftarrow A1,A4} (\\ & \quad \pi_{A1,A4} (\sigma_{A2 = A3 \text{ AND } A4 = C1} (\\ & \quad \quad R1 \bowtie R2 \bowtie \rho_{C1,C2 \leftarrow A1,A2} (R1) \\ & \quad)) \\ &) \end{aligned}$$



SQL: evaluating the queries

- SQL is a declarative language. We are providing the semantics by examples.
- DBMS have *query execution plans* for efficient evaluations:
 - Selections are run as soon as possible
 - When possible, join are ran instead of cartesian products



SQL: formulating the queries

- We don't necessarily have to write efficient queries since DBMS embed query optimizers.
- Hereby it is more important that the provided queries are easy to understand (avoiding errors when formulating the query)



Database Example 2

People

Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Jesse	50	35
Phil	26	30
Louis	50	40
Frank	60	20
Olga	30	41
Steve	85	35
Abby	75	87

Motherhood

Mother	Child
Abby	Alice
Abby	Louis
Jesse	Olga
Jesse	Phil
Alice	Jim
Alice	James

Fatherhood

Father	Child
Steve	Frank
Louis	Olga
Louis	Phil
Frank	Jim
Frank	James



Example 1

- People's fathers earning more than 20

$$\pi_{\text{Father}}(\text{Fatherhood} \bowtie_{\text{Child} = \text{Name}} \sigma_{\text{Income} > 20}(\text{People}))$$

- Same query using SQL:

```
SELECT DISTINCT Father
FROM People, Fatherhood
WHERE Child=Name AND Income > 20
```

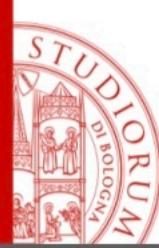


Example 2

- Return the people's name, income and their father's income, where such people earn more than their fathers.

$$\begin{aligned} \pi_{\text{Name}, \text{Income}, \text{IF}} &(\sigma_{\text{Income} > \text{IF}} (\rho_{\text{NF,AF,IF} \leftarrow \text{Name, Age, Income}} (\text{People}) \\ &\quad \bowtie_{\text{NF} = \text{Father}} \\ &\quad (\text{Fatherhood} \bowtie_{\text{Son} = \text{Name}} \text{People})) \\ &) \end{aligned}$$

```
SELECT c.name, c.income, f.income
FROM people f, fatherhood, people c
WHERE f.name = father AND
      child = c.name AND
      c.income > f.income
```



Select with Renaming

```
SELECT c.Name as Name, c.Income AS Income,  
    f.Income AS fatherIncome  
FROM people f, fatherhood, people c  
WHERE f.Name = Father AND Child = c.Name  
    AND c.Income > f.Income
```



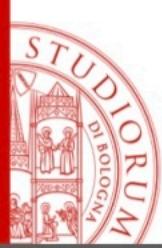
Using Expressions in the Target List

```
SELECT Income/2 AS halvedIncome  
FROM People  
WHERE Name = 'Louis'
```

People

Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Louis	50	40

halvedIncome
20



Join Statement

- Return each person's mother and father

- Implicit JOIN:

```
SELECT f.child, father, mother  
FROM motherhood m, fatherhood f  
WHERE m.child = f.child
```

- Explicit JOIN:

```
SELECT mother, fatherhood.child, father  
FROM motherhood JOIN fatherhood on  
fatherhood.child =  
motherhood.child
```



SELECT with JOIN, Syntax

```
SELECT ...
FROM LeftTable { ... JOIN RightTable ON
  joincondition }, ...
[ WHERE otherPredicate ]
```



Example

- Return name, income and father's income of those people having a greater income than their father's

```
SELECT c.name, c.income, f.income  
FROM people f, fatherhood, people c  
WHERE f.name = father AND  
      child = c.name AND  
      c.income > f.income
```

```
SELECT c.name, c.income, f.income  
FROM (people f JOIN fatherhood ON f.name = father)  
      JOIN people c ON child = c.name  
WHERE c.income > f.income
```



Natural Join

$$\pi_{\text{Child}, \text{Father}, \text{Mother}}(\text{Fatherhood} \bowtie_{\text{Child}=\text{Name}} \rho_{\text{Name} \leftarrow \text{Child}}(\text{Motherhood}))$$

Fatherhood \bowtie Motherhood

```
SELECT mother, Fatherhood.child, father  
FROM MothErhood JOIN Fatherhood ON  
    Fatherhood.child = Motherhood.child
```

```
SELECT mother, child, father  
FROM Motherhood NATURAL JOIN Fatherhood
```



Outer join

- With the previous joins, also called inner joins, some of the tuples could be discarded from the final result: this happens if they don't have a correspondent tuple in the other table.
- In order to avoid this information loss, we can use:

LEFT/RIGHT/FULL OUTER JOIN

- When such join is either left or right, the **OUTER** keyword could be omitted because left and right are “outer” by definition.



Left (outer) join

- Return the father and the mother, if known

```
SELECT Fatherhood.child, father, mother  
FROM Fatherhood LEFT [OUTER] JOIN Motherhood  
ON Fatherhood.child = Motherhood.child
```

Father.Child	Father	Mother
Frank	Steve	NULL
Olga	Louis	Jesse
Phil	Louis	Jesse
Jim	Frank	Alice
James	Frank	Alice



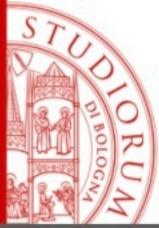
Outer join

```
SELECT Fatherhood.child, father, mother  
FROM Motherhood JOIN Fatherhood  
    ON Motherhood.child = Fatherhood.child
```

```
SELECT Fatherhood.child, father, mother  
FROM Motherhood LEFT OUTER JOIN Fatherhood  
    ON Motherhood.child = Fatherhood.child
```

```
SELECT Fatherhood.child, father, mother  
FROM Motherhood FULL OUTER JOIN Fatherhood  
    ON Motherhood.child = Fatherhood.child
```

- What does the last query return?



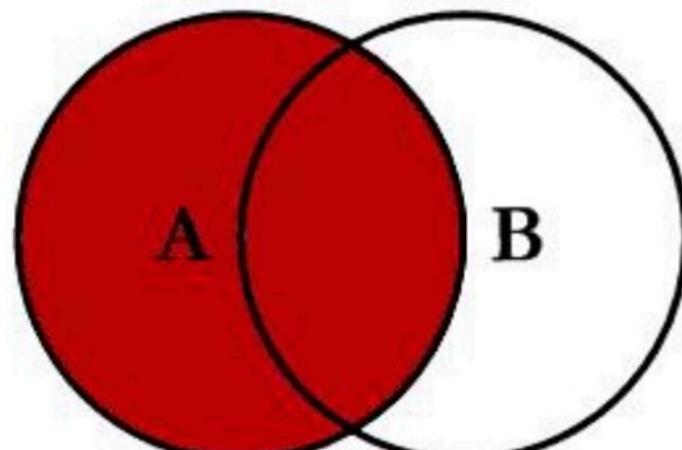
Full Outer join: example

Father.Child	Father	Mother
NULL	NULL	Abby
NULL	NULL	Abby
Olga	Louis	Jesse
Phil	Louis	Jesse
Jim	Frank	Alice
James	Frank	Alice
Frank	Steve	NULL

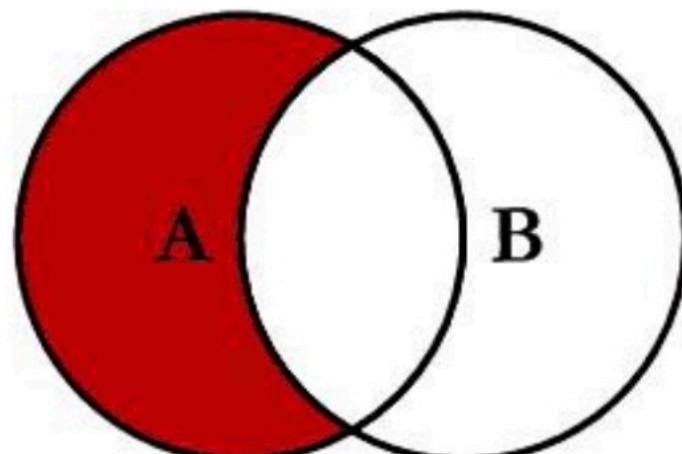
The full outer join returns all the tuples that were excluded on both left and right operand

Recap

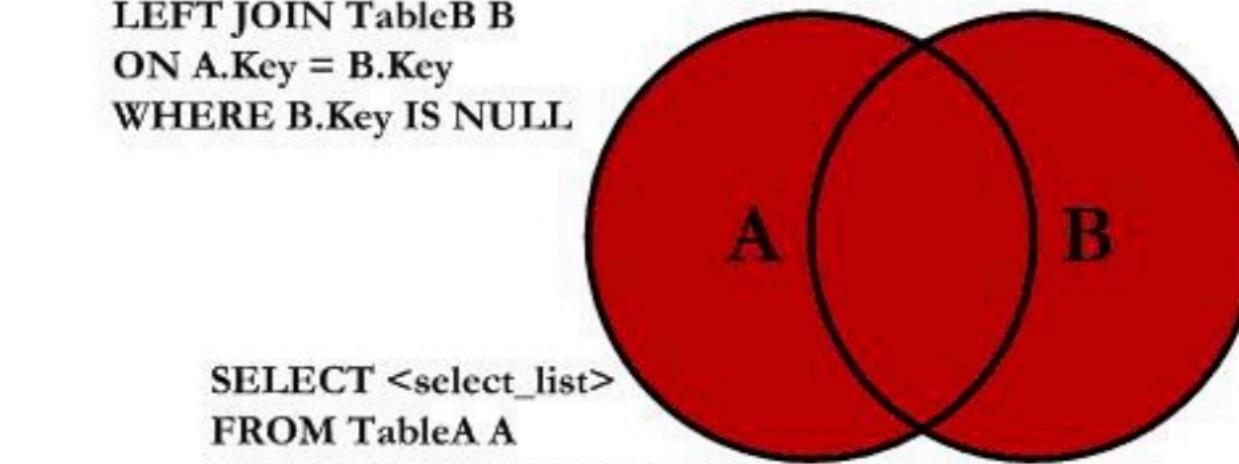
SQL JOINS



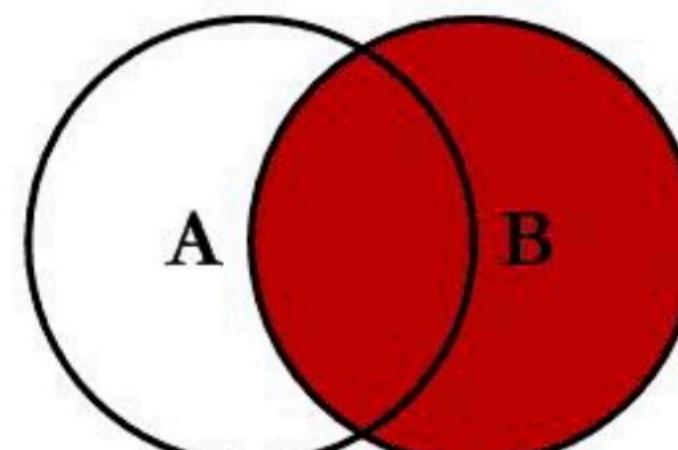
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```



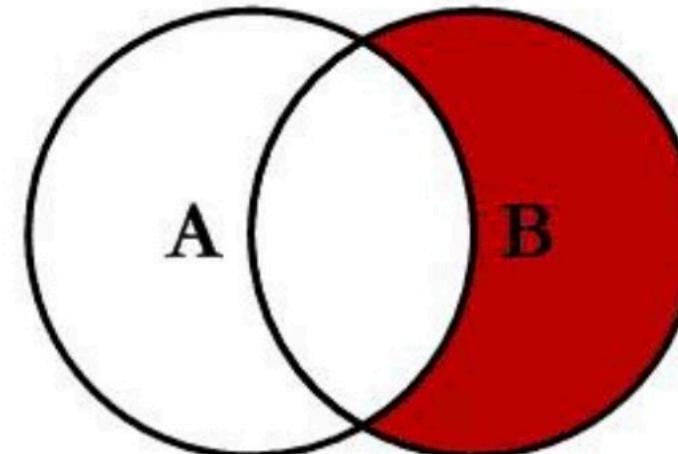
```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```



Sorting the answer

- Provide the name and the income of people being less than 30 yo sorted by **alphabetic order**

```
SELECT name, income  
FROM People  
WHERE age < 30  
ORDER BY name ASC
```

ASC determines an ascending order, while **DESC** a descending one.



Sorting the answer: example

People	Name	Age	Income
	Jim	27	21
	James	25	15
	Alice	55	42
	Jesse	50	35
	Phil	26	30
	Louis	50	40
	Frank	60	20
	Olga	30	41
	Steve	85	35
	Abby	75	87



Sorting the answer

```
SELECT name, income  
FROM People  
WHERE age <= 30
```

```
SELECT name, income  
FROM People  
WHERE age <= 30  
ORDER BY name
```

People

Name	Income
Jim	21
James	15
Phil	30

People

Name	Income
James	15
Jim	21
Phil	30

- **ORDER BY** 's default sorting order is descending



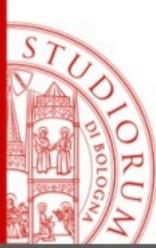
Sorting the answer

```
SELECT name,  
       income  
  FROM People  
 WHERE age <= 30  
 ORDER BY name ASC
```

```
SELECT name,  
       income  
  FROM People  
 WHERE age <= 30  
 ORDER BY name DESC
```

Name	Income
James	15
Jim	21
Phil	30

Name	Income
Phil	30
Jim	21
James	15



Union, Intersection, Difference

- The **SELECT** requires a specific statement for performing unions:

```
SELECT ...
UNION [ALL]
SELECT ...
```

- In the result the rows are unique (except when **all** is used. In this case we have a *multiset union*).

Set Union

```
SELECT child
FROM
Motherhood
UNION
SELECT child
FROM
Fatherhood
```

Motherhood		Fatherhood	
Mother	Child	Father	Child
Abby	Alice	Steve	Frank
Abby	Louis	Louis	Olga
Jesse	Olga	Louis	Phil

Child

Alice
Louis
Olga
Frank
Phil



Multiset Union

```
SELECT child  
FROM  
Motherhood  
UNION ALL  
SELECT child  
FROM  
Fatherhood
```

Motherhood		Fatherhood	
Mother	Child	Father	Child
Abby	Alice	Steve	Frank
Abby	Louis	Louis	Olga
Jesse	Olga	Louis	Phil

Olga is shown
two times

Child
Alice
Louis
Olga
Frank
Olga
Phil



Positional notation (1)

SELECT Father, Child

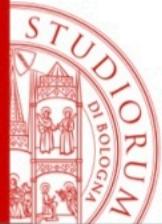
FROM Fatherhood

UNION

SELECT Mother, Child

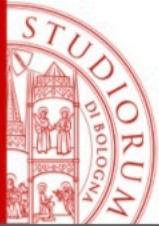
FROM Motherhood

- When two tables have different schema, how could we resolve the conflict by renaming?
 - Either fictitious or none
 - We always assume the first operand's names
 - Merge the conflicting attributes



Positional Notation: first operand

Father	Child
Steve	Frank
Louis	Olga
Louis	Phil
Frank	Jim
Frank	James
Abby	Alice
Abby	Louis
Jesse	Olga
Jesse	Phil
Alice	Jim
Alice	James



Positional notation (2)

```
SELECT father, child  
FROM Fatherhood
```

UNION

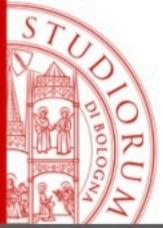
```
SELECT child, mother  
FROM Motherhood
```

```
SELECT father, child  
FROM Fatherhood
```

UNION

```
SELECT mother, child  
FROM Motherhood
```

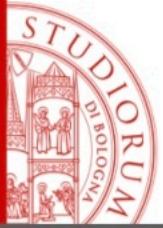
- The resulting tables' resulting scheme in both cases is (father, child)



Difference

```
SELECT Name  
FROM Employee  
EXCEPT  
SELECT Surname AS Name  
FROM Employee
```

- We could later on express such operator through *nested select queries*



Intersection

```
SELECT Name  
FROM Employee  
    INTERSECT  
SELECT Surname AS Name  
FROM Employee
```

is the same as

```
SELECT E.Name  
FROM Employee E, Employee F  
WHERE E.Name = F.Surname
```



Database Example 2

People

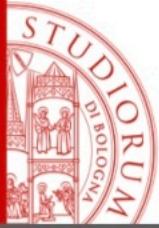
Name	Age	Income
Jim	27	21
James	25	15
Alice	55	42
Jesse	50	35
Phil	26	30
Louis	50	40
Frank	60	20
Olga	30	41
Steve	85	35
Abby	75	87

Motherhood

Mother	Child
Abby	Alice
Abby	Louis
Jesse	Olga
Jesse	Phil
Alice	Jim
Alice	James

Fatherhood

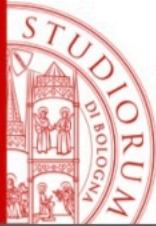
Father	Child
Steve	Frank
Louis	Olga
Louis	Phil
Frank	Jim
Frank	James



Nested Queries

Predicates allow to:

- compare one (or more, as we will see later) attributes with the result of a nested (“sub”) query
- use the existential quantifier (*exists*, \exists)



Nested Queries, example 1

- Provide the name and the income of Frank's father

```
SELECT Name, Income  
FROM People, Fatherhood  
WHERE Name = Father and Child = 'Frank'
```

Here we use the cartesian product and where (equijoin)

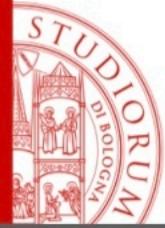
```
SELECT Name, Income  
FROM People  
WHERE Name = (SELECT Father  
         FROM Fatherhood  
         WHERE Child = 'Frank')
```

outer WHERE clause is true when subquery result is equal to Name, and only one tuple is produced by the subquery



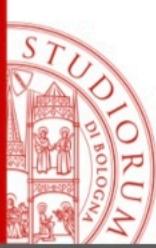
Nested Queries, discussion

- Nested queries are “less declarative”, but sometimes more readable (since they requires less variables)
- Nested and non-nested queried could be combined
- The “subqueries” within nested ones cannot express set operations (“the union can be performed within the *outer query*”); this limitation is not significative.



Semantics of nested queries

- The inner query is performed one time for each tuple within the outer query.



Nested Queries: any, all

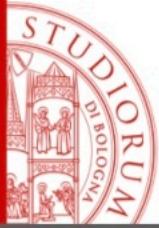
- Nested queries can be formulated through a predicate using either **ANY** or **ALL** alongside with a comparison operator ($>$, $<$, $=$, \geq , ..)

Attribute op **ANY(Expr)**

- An outer query tuple is matched if it satisfies the predicate with respect to any of the tuples within *Expr*

Attribute op **ALL(Expr)**

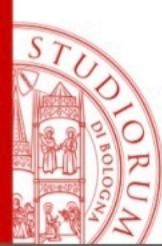
- A outer query tuple is matched if it satisfies the predicate with respect to all the tuples within *Expr*



Nested Queries: IN

Attribute **IN(Expr)**

- An outer query tuple is matched if its values in **Attribute** is contained within the elements returned by *Expr*
- **ANY**, **ALL** and **IN** can be negated through using the word **NOT** before



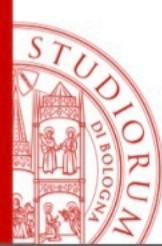
Nested Queries: example 2a

- Provide name and income of the fathers' having child earning more than 20

```
SELECT DISTINCT F.Name, F.Income  
FROM People F, Fatherhood, People C  
WHERE F.Name = Fatherhood.Father AND Fatherhood.Child  
= C.Name AND C.Income > 20
```

We can rewrite it without DISTINCT, because we will not join tables so the fathers' names will not be repeated for each child:

```
SELECT Name, Income  
FROM People  
WHERE Name IN      (SELECT Father  
                      FROM Fatherhood  
                      WHERE Child = ANY (SELECT Name  
                                         FROM People  
                                         WHERE Income > 20))
```



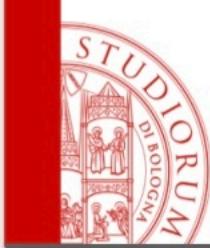
Nested Queries: example 2b

- Provide name and income of the fathers' having child earning more than 20

```
SELECT DISTINCT F.Name, F.Income
FROM People F, Fatherhood, People C
WHERE F.Nome = Father AND Child = C.Name
AND C.Income > 20
```

We can rewrite it without DISTINCT:

```
SELECT Name, Income
FROM People
WHERE Name IN (SELECT Father
FROM Fatherhood, People
WHERE Child=Name AND Income>20)
```



Nested Queries: example 3

- Provide name and income of the fathers' having child earning more than 20, **and provide the child's income too.**

```
SELECT DISTINCT F.Name, F.Income, C.Income  
FROM People F, Fatherhood, People C  
WHERE F.Name = Father AND Child = C.Name  
      AND C.Income > 20
```

- Does the following one provide the same answer?

```
SELECT Name, Income  
FROM People  
WHERE Name IN (SELECT Father  
                FROM Fatherhood  
                WHERE Child = ANY (SELECT Name  
                                    FROM People  
                                    WHERE Income>20))
```

ANY meaning: clause is true if Child value is equal to any of the values returned by nested query

Nested queries visibility

```
SELECT Name, Income
```

```
FROM People
```

```
WHERE Name IN (SELECT Father
```

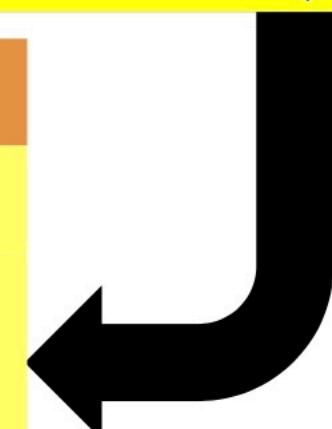
```
FROM Fatherhood
```

```
WHERE Child = ANY (SELECT Name
```

```
FROM People
```

```
WHERE
```

```
Income > 20))
```



Name	Income
Frank	20
Louis	40
Louis	40

Father
Frank
Louis
Louis

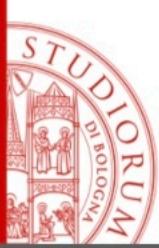
Name
Jim
Alice
Jesse
Phil
Louis
Olga
Steve
Abby

- **Answer:** no, because for each father we do not view the child's income



Nested Queries, considerations

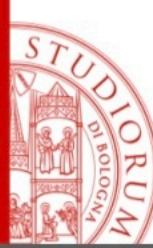
- Visibility Rules:
 - It is not possible to refer to variables declared within inner blocks
 - If a variable's name is omitted, we assume to take the “nearest” declared one
- We can refer to variables declared in outer blocks, but only if they are part of the result of the outer block (e.g. they are among the outer SELECT target list).



Existential Quantification

EXISTS (*Expr*)

- The predicate is true if *Expr* returns at least one tuple



Existential Quantification, example 1

- People having at least one child

```
SELECT *
FROM People
WHERE EXISTS (SELECT *
                    FROM Fatherhood
                    WHERE Father = Name) OR
EXISTS (SELECT *
                    FROM Motherhood
                    WHERE Mother = Name)
```

Name column is taken from the relation obtained in the outer block

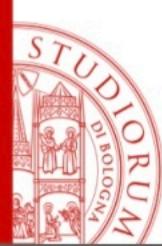


Existential Quantification, example 2

- Fathers whose child earn more than 20

```
SELECT DISTINCT Father
FROM Fatherhood Z
WHERE NOT EXISTS
  (SELECT *
   FROM Fatherhood W, People
   WHERE W.Father = Z.Father
     AND W.child = Name
     AND Income <= 20 )
```

Z.Father is taken from
the relation obtained in
the outer block



Existential Quantification, error

- Fathers having all their children earning more than 20

```
SELECT DISTINCT Father
FROM Fatherhood
WHERE NOT EXISTS ( SELECT *
                     FROM People
                     WHERE Child=Name
                     AND Income<=20 )
```

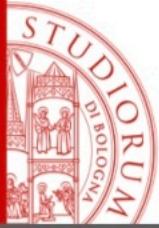
Scope rule: **Child**, without table reference, implicitly refers to the closest **FROM** clause, but **Child** does not belong to table **People**



Existential Quantification, correct

- Fathers having all their children earning more than 20

```
SELECT DISTINCT F.Father
FROM Fatherhood AS F
WHERE NOT EXISTS ( SELECT *
                      FROM People
                      WHERE F.Child=Name
                      AND Income<=20 )
```



Visibility

- Wrong:

```
SELECT *
FROM Employee
WHERE Dept IN ( SELECT Name
                  FROM Department D1
                  WHERE Name = 'Production') OR
      Dept IN ( SELECT Name
                  FROM Department D2
                  WHERE D2.City = D1.City )
```

Wrong because in the last select City of **D1** is not visible



Set difference and Nested Queries

```
SELECT Name FROM Employee  
EXCEPT  
SELECT Surname AS Name FROM Employee
```

```
SELECT Nome  
FROM Employee I  
WHERE NOT EXISTS (SELECT *  
                 FROM Employee  
WHERE Surname = I.Name)
```



Aggregate functions

In the expressions of the target list we can put expressions that compute values from a set of tuples through aggregate functions *Aggr*:

- COUNT, MIN, MAX, AVG, SUM
- Basic syntax:

*Aggr([DISTINCT] *)*

Aggr([DISTINCT] Attribute)



Aggregate Functions: COUNT

- The number of Frank's children

```
SELECT count(*) AS NumFrankChildren
FROM Fatherhood
WHERE father = 'Frank'
```

- the aggregate function (**count**) is applied to the following result's tuples:

```
SELECT *
FROM Fatherhood
WHERE father = 'Frank'
```

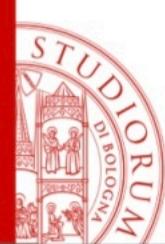


Aggregate Functions: COUNT

Fatherhood

Father	Child
Steve	Frank
Louis	Olga
Louis	Phil
Frank	Jim
Frank	James

NumFrankChildren
2



COUNT DISTINCT

People

Name	Age	Income
Jim	27	30
James	25	24
Alice	55	36
Jesse	50	36

```
SELECT COUNT(*)  
FROM People
```

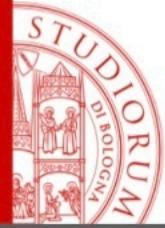
COUNT(*)

4

```
SELECT COUNT(DISTINCT income)  
FROM People
```

COUNT(DISTINC
T income)

3



Some other aggregate functions

- SUM, AVG, MAX, MIN
- Average of the income of Frank's children

```
SELECT AVG(income)
FROM People JOIN Fatherhood ON name=child
WHERE father='Frank'
```



COUNT with null values (1)

People

Name	Age	Income
Jim	27	30
James	25	NULL
Alice	55	36
Jesse	50	36

```
SELECT COUNT(*)  
FROM People
```

count(*)
4

```
SELECT COUNT(income)  
FROM People
```

count(income)
3



COUNT with null values (2)

People

Name	Age	Income
Jim	27	21
James	25	NULL
Alice	55	21
Jesse	50	35

```
SELECT COUNT(DISTINCT income)  
FROM People
```

COUNT(DISTINCT income)
2



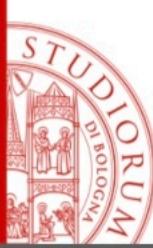
Aggregate Functions and NULLs

People

Name	Age	Income
Jim	27	21
James	25	NULL
Alice	55	21
Jesse	50	35

```
SELECT AVG(income) AS avginc  
FROM People
```

avginc
25,6



Aggregate Functions and Target List

- A wrong query:

```
SELECT name, MAX(income)  
FROM People
```

- Whose the name? We cannot extract the name having the max income. The Target List must have all the same types of attributes.

```
SELECT MIN(age), MAX(income)  
FROM People
```



Maximum and Nested Queries

- Return the people having the (same) maximum income

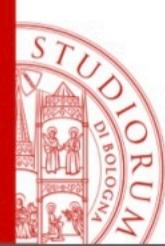
```
SELECT *
FROM People
WHERE income = ( SELECT MAX(income)
                  FROM People )
```



Aggregate Functions and Grouping

- Aggregate function can operate over relations' groups via the **GROUP BY** statement :

GROUP BY AttrList



Aggregate Functions and Grouping

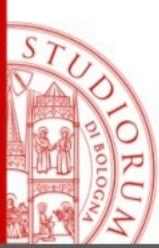
- The number of the fathers' children

```
SELECT Father, COUNT(*) AS NumberofChildren  
FROM Fatherhood  
GROUP BY Father
```

Fatherhood

	Father	Child
	Steve	Frank
	Louis	Olga
	Louis	Phil
	Frank	Jim
	Frank	James

Father	NumberofChildren
Steve	1
Louis	2
Frank	2



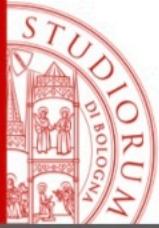
Group By semantics

1. Perform the query without aggregate functions and without aggregate operators

SELECT *

FROM Fatherhood

2. Then perform the grouping and apply the aggregate function over each group



Grouping and Target Lists

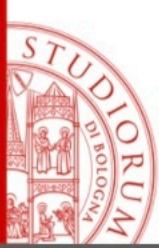
Wrong:

```
SELECT father, AVG(f.income), c.income  
FROM People c JOIN Fatherhood ON child = c.name  
    JOIN People f ON father = f.name  
GROUP BY father
```

we should also group by
c.income because there may
be more than one child

Correct:

```
SELECT father, AVG(f.income), c.income  
FROM People c JOIN Fatherhood ON child = c.name  
    JOIN People f ON father = f.name  
GROUP BY father, c.income
```



Conditions on groups

- Provide those fathers whose children have an average income greater than 25; return the father and their children's average income

```
SELECT father, AVG(f.income)  
FROM People f JOIN Fatherhood ON child = name  
GROUP BY father  
HAVING AVG(f.income) > 25
```



WHERE vs. HAVING

- Provide the fathers whose children under 30 yo have an average income greater than 20

```
SELECT father, AVG(f.income)  
FROM People f JOIN Fatherhood ON child=name  
WHERE age < 30  
GROUP BY father  
HAVING AVG(f.income) > 20
```



Grouping and NULLs

R	A	B
1		11
2		11
3		NULL
4		NULL

**SELECT B, COUNT(*)
FROM R GROUP BY B**

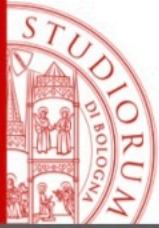
B	COUNT(*)
11	2
NULL	2

**SELECT A, COUNT(*)
FROM R GROUP BY A**

A	COUNT(*)
1	1
2	1
3	1
4	1

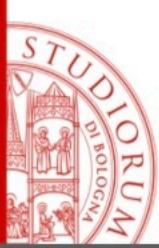
**SELECT A, COUNT(B)
FROM R GROUP BY A**

A	COUNT(B)
1	1
2	1
3	0
4	0



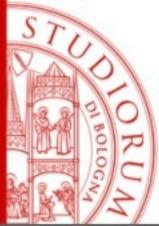
SELECT syntax: summary

```
SELECT AttList1+Exprs  
FROM TableList+Joins  
[ WHERE Condition ]  
[ GROUP BY AttList2]  
[ HAVING AggrCondition]  
[ ORDER BY OrderingAttr1]
```



Updating operations

- Such operations are
 - **INSERT**
 - **DELETE**
 - **UPDATE**
- ...of one or more tuples within a table
- ...on the basis of a predicate that may involve other relations



INSERT

```
INSERT INTO Table [(AttList)]  
VALUES( Vals )
```

or

```
INSERT INTO Table [(AttList)]  
SELECT ...
```



INSERT, examples

```
INSERT INTO People VALUES ('John',25,52)
```

```
INSERT INTO People(Name, Age, Income)  
VALUES('Jack',25,52)
```

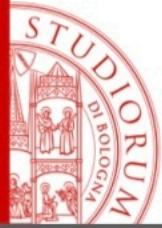
```
INSERT INTO People(Name, Income)  
VALUES('Robert',55)
```

```
INSERT INTO People( Name )  
SELECT Father  
FROM Fatherhood  
WHERE Father NOT IN (SELECT Name  
                      FROM People)
```



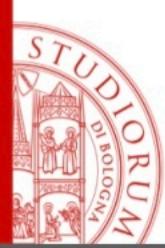
INSERT, discussion

- The attributes' and the values' ordering is relevant
- Both lists should have the same number of arguments
- If the attribute list is omitted, we assume that all the attributes are considered and each value corresponds to a specific attribute as declared in the relation's schema
- If the attribute list does not contain all the relation's attributes, either a NULL value or a default value are emplaced.



Deleting tuples

DELETE FROM Table
[**WHERE** *Condition*]



Deleting tuples, some examples

```
DELETE FROM People  
WHERE Age < 35
```

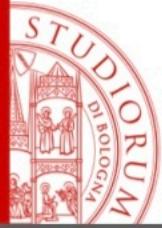
```
DELETE FROM Fatherhood  
WHERE Child NOT IN ( SELECT Name  
                 FROM People)
```

```
DELETE FROM Fatherhood
```



Deleting tuples, discussion

- Removes the tuples satisfying a given condition
- It could cause the removal of other tuples (if the constraints are defined using **CASCADE**)
- If no condition is provided, such has to be intended as **WHERE TRUE**



Updating tuples

```
UPDATE TableName  
SET Attribute = < Expr |  
    SELECT ... |  
    NULL |  
    DEFAULT >  
[ WHERE Condition ]
```



Updating Tuples 1/5

BEFORE

UPDATE People

SET Income = 45

WHERE Name = Bob

AFTER

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	36

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	45



Updating Tuples 2/3

BEFORE

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	36

UPDATE People

SET Income=Income*1.1

WHERE Age < 30

AFTER

People

Name	Age	Income
Jim	27	33
James	25	16,5
Bob	55	36



Updating Tuples 3/5

UPDATE People

```
SET Income =  
  (SELECT Income FROM  
    People WHERE Name=Jim)  
WHERE Name = Bob
```



People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	36

AFTER

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	30



Updating Tuples 4/5

BEFORE

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	36

```
UPDATE People  
SET Income=NULL  
WHERE Age < 30
```

AFTER

People

Name	Age	Income
Jim	27	NULL
James	25	NULL
Bob	55	36



Updating Tuples 5/5

BEFORE

```
UPDATE People  
SET Income=DEFAULT  
WHERE Age < 30
```

AFTER

People

Name	Age	Income
Jim	27	30
James	25	15
Bob	55	36

People

Name	Age	Income
Jim	27	0
James	25	0
Bob	55	36

Assuming that in CREATE TABLE we specified 0 as the DEFAULT value for Income